

POWER8 in-core Cryptography

The Unofficial Guide

Jeffrey Walton
Dr. William Schmidt

POWER8 in-core Cryptography: The Unofficial Guide

by Jeffrey Walton and Dr. William Schmidt

Table of Contents

1. Introduction	1
Organization	1
Compile Farm	1
2. Vector programming	3
PowerPC compilers	3
Altivec headers	3
Machine endianness	3
Malloc and new	4
Vector datatypes	4
Loads and stores	5
3. Advanced Encryption Standard	6
AES encryption	6
AES decryption	6
AES key scheduling	6
4. Secure Hash Standard	7
Sigma functions	7
Ch function	7
Maj function	7
SHA-256	7
SHA-512	7
5. Polynomial multiplication	8
CRC-32 and CRC-32C	8
GCM mode	8
6. Assembly language	9
Cryptogams	9
7. References	10
Cryptogams	10
GitHub	10
IBM	10
NIST	10
8. Revision History	11
Index	12

Chapter 1. Introduction

This document is a guide to using IBM's POWER8 in-core cryptography [<https://www.ibm.com/developerworks/learn/security/index.html>]. The purpose of the book is to document in-core cryptography more completely for developers and quality assurance personnel.

POWER8 in-core cryptography includes CPU instructions to accelerate AES, SHA-256, SHA-512 and polynomial multiplication. This document includes treatments of AES, SHA-256 and SHA-512. It does not include a discussion of polynomial multiplication at the moment, but the chapter is stubbed-out (and waiting for a contributor).

The POWER8 extensions for in-core cryptography find its ancestry in AltiVec SIMD coprocessor. The POWER8 vector unit includes Vector Multimedia Extension (VSX) and the instruction set for in-core cryptography are part it. You can find additional information on VSX at IBM's website at TODO [<https://www.ibm.com/developerworks/>].

Organization

The book proceeds in six parts. First, administrivia is discussed, like how to determine machine endianness and how to load and store a vector from memory. A full treatment of vector programming is its own book, but the discussion should be adequate to move on to the more interesting tasks.

Second, AES is discussed. AES is specified in FIPS 197, Advanced Encryption Standard (AES) [<https://nvlpubs.nist.gov/nistpubs/fips/nist.fips.197.pdf>]. You should read the standard if you are not familiar with the block cipher.

Third, SHA is discussed. SHA is specified in FIPS 180-4, Secure Hash Standard (SHS) [<https://nvlpubs.nist.gov/nistpubs/fips/nist.fips.180-4.pdf>]. You should read the standard if you are not familiar with the hash.

Fourth, polynomial multiplication is discussed. Polynomial multiplications is important for CRC-32, CRC-32C and GCM mode of operation for AES.

Fifth, performance is discussed. The implementations are compared against C and C++ routines and assembly language routines from OpenSSL. The OpenSSL routines are high quality and written by Andy Polyakov.

Finally, assembly language integration is discussed. Andy Polyakov dual licenses his cryptographic implementations and you can use his routines once you know how to integrate them.

Compile Farm

If you are a free and open software developer then you are eligible for a free GCC Compile Farm [<https://cfarm.tetaneutral.net/>] account. The Compile Farm provides machines for different architectures, including MIPS64, Aarch64 and PowerPC64. Access is provided through SSH.

The Compile Farm offers four 64-bit PowerPC machines, and two of them can be used for POWER8 testing. The two machines are `gcc112.fsffrance.org` and `gcc119.fsffrance.org`. The first machine is `gcc112`, and it is a Linux `ppc64-le` machine (PowerPC, 64-bit, little-endian). The second machine is `gcc119`, and it is a AIX `ppc64-be` machine (PowerPC, 64-bit, big-endian). Both machines are IBM POWER System S822 with two CPU cards. `gcc112` has 160 logical CPUs, and `gcc119` has 64 logical CPUs.

TODO: This book is free software. If you see an opportunity for improvement, and error or an omission then submit a pull request or open a bug report.

Chapter 2. Vector programming

AES, SHA-256 and SHA-512 acceleration occurs in the VSX unit. Data must be moved from main memory into a vector register, the data must be transformed, and then data must be written back to main memory. Data is moved to and from main memory using vector loads and stores.

Several topics need to be discussed leading to trouble free loads and stores. They include PowerPC compilers and options, AltiVec headers, machine endianness, vector datatypes and then finally the loads and stores.

PowerPC compilers

This documents uses two compilers for testing. The first is GCC and the second is IBM XL C/C++. Each compiler is slightly different with its options.

Compiling a test program with GCC will generally look like below. The important part is `-mcpu=power8` which selects the POWER8 Instruction Set Architecture (ISA).

```
$ g++ -mcpu=power8 test.cxx -o test.exe
```

Complimentary, compiling a test program with IBM XL C/C++ will generally look like below. The important parts are the C++ compiler name of `xlc`, and `-qarch=pwr8` which selects the POWER8 ISA.

```
$ xlc -qarch=pwr8 -qaltivec test.cxx -o test.exe
```

When compiling source code to examine the quality of code generation the program should be compiled with `-O3`. Both compilers consume `-O3`.

AltiVec headers

The header required for datatypes and functions is `<altivec.h>`. To support compilers with a C++ compiler `__vector` keyword is used rather than `vector`. A typical AltiVec include looks as shown below.

```
#if defined(__ALTIVEC__)
# include <altivec.h>
# undef vector
# undef pixel
# undef bool
#endif
```

Machine endianness

You will experience both little-endian and big-endian machines in the field when working with a modern PowerPC architecture. Linux is generally little-endian, while AIX is big-endian. When writing portable machine code you should check the value of preprocessor macros `__LITTLE_ENDIAN__` or `__BIG_ENDIAN__`.

The value of the macros `__BIG_ENDIAN__` and `__LITTLE_ENDIAN__` are defined to non-0 to indicate endianness. Your source code should look similar to shown below.

```
#if __LITTLE_ENDIAN__
# error "Little-endian system"
```

```
#else
# error "Big-endian system"
#endif
```

If you are in doubt about the endianness preprocessor macros you can ask the compiler to tell what they are. Below is from GCC on gcc112, which is ppc64-le on the compile farm.

```
$ g++ -dM -E test.cxx | grep -i endian
#define __ORDER_LITTLE_ENDIAN__ 1234
#define _LITTLE_ENDIAN 1
#define __FLOAT_WORD_ORDER__ __ORDER_LITTLE_ENDIAN__
#define __ORDER_PDP_ENDIAN__ 3412
#define __LITTLE_ENDIAN__ 1
#define __ORDER_BIG_ENDIAN__ 4321
#define __BYTE_ORDER__ __ORDER_LITTLE_ENDIAN__
```

And the complimentary view from IBM XL C/C++ on gcc112, which is ppc64-le on the compile farm.

```
$ xlc -qshowmacros -E test.cxx | grep -i endian
#define _LITTLE_ENDIAN 1
#define __BYTE_ORDER__ __ORDER_LITTLE_ENDIAN__
#define __FLOAT_WORD_ORDER__ __ORDER_LITTLE_ENDIAN__
#define __LITTLE_ENDIAN__ 1
#define __ORDER_BIG_ENDIAN__ 4321
#define __ORDER_LITTLE_ENDIAN__ 1234
#define __ORDER_PDP_ENDIAN__ 3412
#define __VEC_ELEMENT_REG_ORDER__ __ORDER_LITTLE_ENDIAN__
```

However, below is gcc119 from the compile farm, which is ppc64-be. It runs AIX and notice `__BYTE_ORDER__`, `__ORDER_BIG_ENDIAN__` and `__ORDER_LITTLE_ENDIAN__` are not defined.

```
$ xlc -qshowmacros -E test.cxx | grep -i endian
#define __BIG_ENDIAN__ 1
#define _BIG_ENDIAN 1
#define __THW_BIG_ENDIAN__ 1
#define __HHW_BIG_ENDIAN__ 1
```

Malloc and new

The system calls `malloc` and `new` (and friends) are used to acquire memory from the heap. The system calls *do not* guarantee alignment to any particular boundary on all platforms. Linux generally returns a pointer that is at least 16-byte aligned on all platforms, including ARM, PPC, MIPS and x86. AIX does not provide the same alignment behavior.

To avoid unexpected surprises when using heap allocations you should use `posix_memalign` to acquire heap memory aligned to a particular address.

TODO: I believe AIX or XLC has another function to call for vector programming.

Vector datatypes

Three vector datatypes are used for in-core programming. They are .

Loads and stores

Altivec loads and stores have traditionally been performed using `vec_ld` and `vec_st` since at least the POWER4 days in the 1990s. `vec_ld` and `vec_st` are sensitive to alignment of the effective memory address. The effective address is the `address+offset` rounded down or masked to a multiple of 16.

The effective address passed to `vec_ld` and `vec_st` must be aligned to a 16-byte boundary or incorrect results will arise. Altivec *does not* raise a SIGBUS. Instead, the bottom 4 bits of the address are masked-off and the value at the effective address is read.

POWER7 introduced unaligned loads and stores that avoid the aligned memory address requirement. The instructions to use for unaligned loads and stores are `vec_vsx_ld` and `vec_vsx_st` when using GCC; and `vec_xl` and `vec_xst` when using XLC.

Chapter 3. Advanced Encryption Standard

AES is the Advanced Encryption Standard. AES is specified in FIPS 197, Advanced Encryption Standard (AES) [<https://nvlpubs.nist.gov/nistpubs/fips/nist.fips.197.pdf>]. You should read the standard if you are not familiar with the block cipher.

Three topics are discussed for AES. The first is encryption, the second is decryption, and the third is keying. Keying is discussed last because encryption and decryption uses the golden key schedule from FIPS 197.

AES encryption

XXX

AES decryption

XXX

AES key scheduling

XXX

Chapter 4. Secure Hash Standard

SHA is the Secure Hash Standard. SHA is specified in FIPS 180-4, Secure Hash Standard (SHS) [<https://nvlpubs.nist.gov/nistpubs/fips/nist.fips.180-4.pdf>]. You should read the standard if you are not familiar with the hash.

Sigma functions

POWER8 provides a `sigma` instruction to accelerate SHA calculations. The instruction takes two integer arguments and the constants are used to select among `Sigma0`, `Sigma1`, `sigma0` and `sigma1`.

Ch function

POWER8 provides the `vsel` instruction and it is SHA's Ch function. The implementation for the 32x4 arrangement is shown below. The code is the same for the 64x2 arrangement, but the function takes `uint64x2_p8` arguments. The important piece of information is `x` is used as the selector.

```
uint32x4_p8 inline
VectorCh(const uint32x4_p8 x, const uint32x4_p8 y, const uint32x4_p8 z)
{
    return vec_sel(z, y, x);
}
```

Maj function

POWER8 provides the `vsel` instruction and it can be used for SHA's Maj function. The implementation for the 32x4 arrangement is shown below. The code is the same for the 64x2 arrangement, but the function takes `uint64x2_p8` arguments. The important piece of information is `x^y` is used as the selector.

```
uint32x4_p8 inline
VectorCh(const uint32x4_p8 x, const uint32x4_p8 y, const uint32x4_p8 z)
{
    return vec_sel(y, z, vec_xor(x, y));
}
```

SHA-256

XXX

SHA-512

XXX

Chapter 5. Polynomial multiplication

CRC-32 and CRC-32C

xxx

GCM mode

xxx

Chapter 6. Assembly language

XXX.

Cryptogams

Cryptogams [<https://www.openssl.org/~appro/cryptogams/>] is Andy Polyakov's incubator to develop assembly language routines for OpenSSL. Andy dual licenses his implementations, so a more permissive license is available for the assembly language source code. This chapter will show you how to build Andy's software.

Chapter 7. References

The following is a list of references we are aware of.

Cryptogams

The following is a list of Cryptogams references.

- CRYPTOGRAMS: low-level cryptographic primitives collection [<https://www.openssl.org/~appro/cryptogams/>]

GitHub

The following is a list of GitHub references.

- AES Intrinsic [<https://github.com/nolader/AES-Intrinsic>]
- SHA Intrinsic [<https://github.com/nolader/SHA-Intrinsic>]
- CRC32/vpmsum [<https://github.com/antonblanchard/crc32-vpmsum>]

IBM

The following is a list of IBM references.

- Recommended debug, compiler, and linker settings for Power processor tuning [<https://www.ibm.com/support/knowledgecenter/en/linuxonibm/liaal/iplsdkrecbldset.htm>]
- POWER8 in-core cryptography [<https://www.ibm.com/developerworks/library/se-power8-in-core-cryptography/index.html>]

NIST

The following is a list of NIST references.

- FIPS 197, Advanced Encryption Standard (AES) [<https://nvlpubs.nist.gov/nistpubs/fips/nist.fips.197.pdf>]
- FIPS 180-4, Secure Hash Standard (SHS) [<https://nvlpubs.nist.gov/nistpubs/fips/nist.fips.180-4.pdf>]

Chapter 8. Revision History

Revision History

Revision 1

1 April 2018

jww, ws

Initial release

Index

C

Cryptogams, 9

S

SHA, 7

Ch function, 7

Maj function, 7

SHA-256, 7

SHA-512, 7

Sigma functions, 7