

POWER8 in-core Cryptography

The Unofficial Guide

Jeffrey Walton
Dr. William Schmidt

POWER8 in-core Cryptography: The Unofficial Guide

by Jeffrey Walton and Dr. William Schmidt

Extensive review and rough drafts: Segher Boessenkool

Publication date 1 April 2018

Table of Contents

1. Introduction	1
Organization	1
Compile Farm	1
2. Vector programming	3
PowerPC compilers	3
Altivec headers	3
Machine endianness	4
Malloc and new	4
Vector datatypes	5
Vector shifts	5
Aligned loads	6
Unaligned loads	7
Big-endian loads	7
Const pointers	8
3. Runtime features	9
AIX features	9
Linux features	9
SIGILL probes	10
L1 Data Cache	10
4. Advanced Encryption Standard	11
AES encryption	11
AES decryption	11
AES key scheduling	11
5. Secure Hash Standard	12
Sigma functions	12
Ch function	12
Maj function	12
SHA-256	12
SHA-512	12
6. Polynomial multiplication	13
CRC-32 and CRC-32C	13
GCM mode	13
7. Assembly language	14
Cryptogams	14
8. References	15
Cryptogams	15
GitHub	15
IBM	15
NIST	15
Stack Exchange	15
Index	16

Chapter 1. Introduction

This document is a guide to using IBM's POWER8 in-core cryptography [<https://www.ibm.com/developerworks/learn/security/index.html>]. The purpose of the book is to document in-core cryptography more completely for developers and quality assurance personnel who wish to take advantage of the features.

POWER8 in-core cryptography includes CPU instructions to accelerate AES, SHA-256, SHA-512 and polynomial multiplication. This document includes treatments of AES, SHA-256 and SHA-512. It does not include a discussion of polynomial multiplication at the moment, but the chapter is stubbed-out (and waiting for a contributor).

The POWER8 extensions for in-core cryptography find its ancestry in AltiVec SIMD coprocessor. The POWER8 vector unit includes Vector Multimedia Extension (VSX) and the instruction set for in-core cryptography are part it. You can find additional information on VSX at IBM's website at TODO [<https://www.ibm.com/developerworks/>].

The source code in the book is a mix of C and C++. The SHA-256 and SHA-512 samples were written in C++ to avoid compile errors due to the SHA API requiring 4-bit literal constants. We could not pass parameters through functions and obtain the necessary `constexpr` so template parameters were used instead.

Organization

The book proceeds in six parts. First, administration is discussed, like how to determine machine endianness and how to load and store a vector from memory. A full treatment of vector programming is its own book, but the discussion should be adequate to move on to the more interesting tasks.

Second, AES is discussed. AES is specified in FIPS 197, Advanced Encryption Standard (AES) [<https://nvlpubs.nist.gov/nistpubs/fips/nist.fips.197.pdf>]. You should read the standard if you are not familiar with the block cipher.

Third, SHA is discussed. SHA is specified in FIPS 180-4, Secure Hash Standard (SHS) [<https://nvlpubs.nist.gov/nistpubs/fips/nist.fips.180-4.pdf>]. You should read the standard if you are not familiar with the hash.

Fourth, polynomial multiplication is discussed. Polynomial multiplication is important for CRC-32, CRC-32C and GCM mode of operation for AES.

Fifth, performance is discussed. The implementations are compared against C and C++ routines and assembly language routines from OpenSSL. The OpenSSL routines are high quality and written by Andy Polyakov.

Finally, assembly language integration is discussed. Andy Polyakov dual licenses his cryptographic implementations and you can use his routines once you know how to integrate them.

Compile Farm

The book makes frequent references to `gcc112` and `gcc119` from the GCC Compile Farm. The Compile Farm offers four 64-bit PowerPC machines, and `gcc112` and `gcc119` are the POWER8 iron (the other two are POWER7 hardware). `gcc112` is a Linux PowerPC, 64-bit, little-endian machine (`ppc64-le`), and `gcc119` is an AIX PowerPC, 64-bit, big-endian machine (`ppc64-be`).

Both POWER8 machines are IBM POWER System S822 with two CPU cards. `gcc112` has 160 logical CPUs and runs at 3.4 GHz. `gcc119` has 64 logical CPUs and runs at 4.1 GHz. At 4.1 GHz and 192 GB of RAM `gcc119` is probably a contender for one of the fastest machine you will work on.

If you are a free and open software developer then you are eligible for a free GCC Compile Farm [<https://cfarm.tetaneutral.net/>] account. The Cfarm provides machines for different architectures, including MIPS64, Aarch64 and PowerPC64. Access is provided through SSH.

TODO: This book is free software. If you see an opportunity for improvement, and error or an omission then submit a pull request or open a bug report.

Chapter 2. Vector programming

Several topics need to be discussed to minimize trouble when using the AltiVec and POWER8 extensions. They include PowerPC compilers and options, AltiVec headers, machine endianness, vector datatypes, memory and alignment, and loads and stores. It is enough information to get to the point you can use AES and SHA but not much more.

Memory alignment, loads, stores and shifts will probably cause the most trouble for someone new to PowerPC vector programming. If you are new to the platform you may want to read this chapter twice. If you are experienced with the platform then you probably want to skip this chapter.

PowerPC compilers

Two compilers are used for testing. The first is GCC and the second is IBM XL C/C++. The compilers are mostly the same but slightly different with its options.

Compiling a test program with GCC will generally look like below. The important part is `-mcpu=power8` which selects the POWER8 Instruction Set Architecture (ISA).

```
$ g++ -mcpu=power8 test.cxx -o test.exe
```

Complimentary, compiling a test program with IBM XL C/C++ will generally look like below. The important parts are the C++ compiler name of `xlc`, and `-qarch=pwr8` which selects the POWER8 ISA.

```
$ xlc -qarch=pwr8 -qaltivec test.cxx -o test.exe
```

When compiling source code to examine the quality of code generation the program should be compiled with `-O3`. Both compilers consume `-O3`.

AltiVec headers

The header required for datatypes and functions is `<altivec.h>`. To support compiles with a C++ compiler `__vector` keyword is used rather than `vector`. A typical AltiVec include looks as shown below.

```
#if defined(__ALTIVEC__)
# include <altivec.h>
# undef vector
# undef pixel
# undef bool
#endif
```

In addition to `__ALTIVEC__` preprocessor macro you will see the following defines depending on the platform:

- `__powerpc__` and `__powerpc` on AIX
- `__powerpc__` and `__powerpc64__` on Linux
- `_ARCH_PWR3` through `_ARCH_PWR9` on AIX and Linux
- `__linux__`, `__linux` and `linux` on Linux
- `_AIX`, and `_AIX32` through `_AIX72` on AIX
- `__xlc__` and `__xlc` when using IBM XL C/C++

Machine endianness

You will experience both little-endian and big-endian machines in the field when working with a modern PowerPC architecture. Linux is generally little-endian, while AIX is big-endian.

When writing portable source code you should check the value of preprocessor macros `__LITTLE_ENDIAN__` or `__BIG_ENDIAN__` to determine the configuration. The value of the macros `__BIG_ENDIAN__` and `__LITTLE_ENDIAN__` are defined to non-0 to activate the macro. Source code checking endianness should look similar to the code shown below.

```
#if __LITTLE_ENDIAN__
# error "Little-endian system"
#else
# error "Big-endian system"
#endif
```

The compilers can show the endian related preprocessor macros available on a platform. Below is from GCC on gcc112 from the compile farm, which is ppc64-le.

```
$ g++ -dM -E test.cxx | grep -i endian
#define __ORDER_LITTLE_ENDIAN__ 1234
#define __LITTLE_ENDIAN__ 1
#define __FLOAT_WORD_ORDER__ __ORDER_LITTLE_ENDIAN__
#define __ORDER_PDP_ENDIAN__ 3412
#define __LITTLE_ENDIAN__ 1
#define __ORDER_BIG_ENDIAN__ 4321
#define __BYTE_ORDER__ __ORDER_LITTLE_ENDIAN__
```

And the complimentary view from IBM XL C/C++ on gcc112 from the compile farm, which is ppc64-le.

```
$ xlc -qshowmacros -E test.cxx | grep -i endian
#define __LITTLE_ENDIAN__ 1
#define __BYTE_ORDER__ __ORDER_LITTLE_ENDIAN__
#define __FLOAT_WORD_ORDER__ __ORDER_LITTLE_ENDIAN__
#define __LITTLE_ENDIAN__ 1
#define __ORDER_BIG_ENDIAN__ 4321
#define __ORDER_LITTLE_ENDIAN__ 1234
#define __ORDER_PDP_ENDIAN__ 3412
#define __VEC_ELEMENT_REG_ORDER__ __ORDER_LITTLE_ENDIAN__
```

However, below is gcc119 from the compile farm, which is ppc64-be. It runs AIX and notice `__BYTE_ORDER__`, `__ORDER_BIG_ENDIAN__` and `__ORDER_LITTLE_ENDIAN__` are not present.

```
$ xlc -qshowmacros -E test.cxx | grep -i endian
#define __BIG_ENDIAN__ 1
#define __BIG_ENDIAN__ 1
#define __THW_BIG_ENDIAN__ 1
#define __HWP_BIG_ENDIAN__ 1
```

Malloc and new

The system calls `malloc` and `new` (and friends) are used to acquire memory from the heap. The system calls *do not* guarantee alignment to any particular boundary on all platforms. Linux generally returns a

pointer that is at least 16-byte aligned on all platforms, including ARM, PPC, MIPS and x86. AIX does not provide the same alignment behavior.

To avoid unexpected surprises when using heap allocations you should use `posix_memalign` to acquire heap memory aligned to a particular boundary.

TODO: I believe AIX or XLC has another function to call for memory used with vector programming.

Vector datatypes

Three vector datatypes are needed for in-core programming. The three types used for crypto are listed below.

- `__vector unsigned char`
- `__vector unsigned int`
- `__vector unsigned long`

`__vector unsigned char` is arranged as 16 each 8-bit bytes, and it is typedef'd as `uint8x16_p8`.
`__vector unsigned int` is arranged as 4 each 32-bit words, and it is typedef'd as `uint32x4_p8`.

`__vector unsigned long` is new to POWER8. `__vector unsigned long` is arranged as 2 each 64-bit double words, and it is typedef'd as `uint64x2_p8`.

The typedef naming was selected to convey the arrangement, like 32x4 and 64x2. The trailing `_p8` was selected to avoid collisions with ARM NEON vector data types. The suffix `_p` (for POWER architecture) or `_v` (for Vector) would work just as well.

Vector shifts

Altivec shifts and rotates are performed using *Vector Shift Left Double by Octet Immediate*. The vector shift and rotate built-in is `vec_sld` and it compiles/assembles to `vsldoi`. Both shift and rotate operate on a concatenation of two vectors. Bytes are shifted out on the left and shifted in on the right. The instructions need an integral constant in the range 0 - 15, inclusive.

Vector shifts and rotates perform as expected on big-endian machines. Little-endian machines need a special handling to produce correct results and the IBM manuals don't tell you about it [https://www.ibm.com/support/knowledgecenter/SSXVZZ_13.1.4/com.ibm.xlcpp1314.linux.doc/compiler_ref/vec_sld.html]. If you are like many other developers then you will literally waste hours trying to figure it out what happened the first time you experience it.

The issue is shifts and rotates are endian sensitive [<https://stackoverflow.com/q/46341923/608639>], and you have to use `16-n` and swap vector arguments on little-endian systems. The C++ source code provides the following template function to compensate for the little-endian behavior.

```
template <unsigned int N, class T>
T VectorShiftLeft(const T val1, const T val2)
{
    #if __LITTLE_ENDIAN__
        enum {R = (16-N)&0xf};
        return vec_sld(val2, val1, R);
    #else
```



```
        enum {R = N&0xf};  
        return vec_sld(val1, val2, R);  
#endif  
}
```

A `VectorRotateLeft` would be similar to the code below, if needed. Rotate is a special case of shift where both vector arguments are the same value.

```
template <unsigned int N, class T>  
T VectorRotateLeft(const T val)  
{  
#if __LITTLE_ENDIAN__  
    enum {R = (16-N)&0xf};  
    return vec_sld(val, val, R);  
#else  
    enum {R = N&0xf};  
    return vec_sld(val, val, R);  
#endif  
}
```

Aligned loads

Altivec loads and stores have traditionally been performed using `vec_ld` and `vec_st` since at least the POWER4 days in the early 2000s. `vec_ld` and `vec_st` are sensitive to alignment of the memory address and the offset into the address. The effective address is the address+offset rounded down or masked to a multiple of 16.

The effective address used for `vec_ld` and `vec_st` must be aligned to a 16-byte boundary or incorrect results will arise. Altivec *does not* raise a SIGBUS. Instead, the bottom 4 bits of the address are masked-off and then the memory at the effective address is loaded.

You can use the Altivec loads and stores whenever you *control* buffers and their alignments, like AES key schedule table. Otherwise just use unaligned loads and stores to avoid trouble.

The C/C++ code to perform a load using `vec_ld` should look similar to below. Notice the `assert` to warn you of problems in debug builds.

```
template <class T>  
uint32x4_p8 VectorLoad(const T* mem_addr, int offset)  
{  
#ifndef NDEBUG  
    uintptr_t maddr = ((uintptr_t)mem_addr)+offset;  
    uintptr_t mask = ~(uintptr_t)0xf;  
    uintptr_t eaddr = maddr & mask;  
    assert(maddr == eaddr);  
#endif  
  
    return (uint32x4_p8)vec_ld(offset, (uint8_t*)mem_addr);  
}
```

The C/C++ code to perform a store using `vec_st` should look similar to below.

```
template <class T>  
void VectorStore(T* mem_addr, int offset)
```

```
{
#ifdef NDEBUG
    uintptr_t maddr = ((uintptr_t)mem_addr)+offset;
    uintptr_t mask = ~(uintptr_t)0xf;
    uintptr_t eaddr = maddr & mask;
    assert(maddr == eaddr);
#endif

    vec_st(offset, (uint8_t*)mem_addr);
}
```

Casting away the const-ness on mem_addr is discussed in const pointers below.

Unaligned loads

POWER7 introduced unaligned loads and stores that avoid the aligned memory requirements. The instructions for unaligned loads and stores are `vec_vsx_ld` and `vec_vsx_st` for GCC; and `vec_xl` and `vec_xst` for XLC.

You should use the POWER7 loads and stores whenever you *do not* control buffers or their alignments, like messages supplied by user code.

The C/C++ code to perform a load using `vec_xl` and `vec_vsx_ld` should look similar to below. An `assert` is no longer needed for debug builds, and the function name has a `u` added to indicate unaligned.

```
template <class T>
uint32x4_p8 VectorLoadu(const T* mem_addr, int offset)
{
#ifdef __xlc__ || defined(__xlC__)
    return (uint32x4_p8)vec_xl(offset, (uint8_t*)mem_addr);
#else
    return (uint32x4_p8)vec_vsx_ld(offset, (uint8_t*)mem_addr);
#endif
}
```

The C/C++ code to perform a store using `vec_xst` and `vec_vsx_st` should look similar to below.

```
template <class T>
void VectorStoreu(T* mem_addr, int offset)
{
#ifdef __xlc__ || defined(__xlC__)
    vec_xst((uint8x16_p8)val, offset, (uint8_t*)mem_addr);
#else
    vec_vsx_st((uint8x16_p8)val, offset, (uint8_t*)mem_addr);
#endif
}
```

Casting away the const-ness on mem_addr is discussed in const pointers below.

Big-endian loads

TODO: talk about `vec_xl_be` when using XLC, and the assembly replacement `VEC_XL_BE` when using GCC.

Const pointers

The AltiVec built-ins have an unusual behavior when it comes to `const` pointers used during a load operation. A program runs slower when the memory is marked as `const`. The behavior has been witnessed in two libraries and the decrease in performance is measurable. For example, AES runs 0.3 to 0.5 cycles per byte (cpb) faster [<https://github.com/randombit/botan/pull/1459>] when `non-const` pointers are used for loads. 0.3 cpb may not sound like much but it equates to 200 MiB/s for AES-128 on `gcc112`.

Source code using the `non-const` pointers should look similar to below:

```
template <class T>
uint32x4_p8 VectorLoad(const T* mem_addr, int offset)
{
    // mem_addr must be aligned to 16-byte boundary
    return (uint32x4_p8)vec_ld(offset, (uint8_t*)mem_addr);
}
```

Chapter 3. Runtime features

Runtime feature detections allows code to switch to a faster implementation when the hardware permits. This chapter shows you how to determine in-core crypto availability at runtime on AIX and Linux PowerPC platforms.

AIX features

TODO: find out how to perform runtime feature detection on AIX. We checked `getsystemcfg` and `sysconf` for ISA 2.07, polynomial multiply, AES and SHA (and crypto) bits but they are missing.

The only thing we have found is `SIGILL` probes and signal handlers. It would be nice to avoid the nastiness.

Linux features

Some versions of Glibc and the kernel provide ELF auxiliary vectors with the information. `AT_HWCAP2` will show the `vcrypto` flag when in-core crypto is available. TODO: which versions?

```
$ LD_SHOW_AUXV=1 /bin/true
AT_DCACHEBSIZE: 0x80
AT_ICACHEBSIZE: 0x80
AT_UCACHEBSIZE: 0x0
AT_SYSINFO_EHDR: 0x3fff877c0000
AT_HWCAP:      ppcle true_le archpmu vsx arch_2_06 dfp ic_snoop
               smt mmu fpu altivec ppc64 ppc32
AT_PAGESZ:     65536
AT_CLKTCK:     100
AT_PHDR:       0x10000040
AT_PHEENT:     56
AT_PHNUM:      9
AT_BASE:       0x3fff877e0000
AT_FLAGS:      0x0
AT_ENTRY:      0x1000145c
AT_UID:        10455
AT_EUID:       10455
AT_GID:        10455
AT_EGID:       10455
AT_SECURE:     0
AT_RANDOM:     0x3fffeaeaa872
AT_HWCAP2:     vcrypto tar isel ebb dscr htm arch_2_07
AT_EXECFN:     /bin/true
AT_PLATFORM:   power8
AT_BASE_PLATFORM: power8
```

Linux systems with Glibc version 2.16 can use `getauxval` to determine CPU features. Runtime code to perform the check should look similar to below. The defines were taken from the Linux kernel's `cputable.h` [<https://git.kernel.org/pub/scm/linux/kernel/git/torvalds/linux.git/tree/arch/powerpc/include/asm/cputable.h>].

```
#ifndef AT_HWCAP2
# define AT_HWCAP2 26
```

```
#endif
#ifndef PPC_FEATURE2_ARCH_2_07
# define PPC_FEATURE2_ARCH_2_07    0x80000000
#endif
#ifndef PPC_FEATURE2_VEC_CRYPTO
# define PPC_FEATURE2_VEC_CRYPTO    0x02000000
#endif

bool HasPower8()
{
    if (getauxval(AT_HWCAP2) & PPC_FEATURE2_ARCH_2_07 != 0)
        return true;
    return false;
}

bool HasPower8Crypto()
{
    if (getauxval(AT_HWCAP2) & PPC_FEATURE2_VEC_CRYPTO != 0)
        return true;
    return false;
}
```

SIGILL probes

TODO: show this nasty technique.

L1 Data Cache

The L1 data cache line size is an important security parameter that can be used to avoid leaking information through timing attacks. IBM POWER System S822, like gcc112 and gcc119, have a 128-byte L1 data cache line size.

gcc119 runs AIX and a program can query the L1 data cache line size as shown below.

```
#include <sys/systemcfg.h>

int cacheLineSize = getsystemcfg(SC_L1C_DLS);
if (cacheLineSize) <= 0)
    cacheLineSize = DEFAULT_L1_CACHE_LINE_SIZE;
```

gcc112 runs Linux and a program can query the L1 data cache line size as shown below.

```
#include <sys/sysconf.h>

int cacheLineSize = sysconf(_SC_LEVEL1_DCACHE_LINESIZE);
if (cacheLineSize) <= 0)
    cacheLineSize = DEFAULT_L1_CACHE_LINE_SIZE;
```

It is important to check the return value from sysconf on Linux. gcc112 runs CentOS 7.4 and the machine returns 0 for the L1 cache line query. Also see sysconf and _SC_LEVEL1_DCACHE_LINESIZE returns 0? [<https://lists.centos.org/pipermail/centos/2017-September/166236.html>] on the CentOS mailing list.

Chapter 4. Advanced Encryption Standard

AES is the Advanced Encryption Standard. AES is specified in FIPS 197, Advanced Encryption Standard (AES) [<https://nvlpubs.nist.gov/nistpubs/fips/nist.fips.197.pdf>]. You should read the standard if you are not familiar with the block cipher.

Three topics are discussed for AES. The first is encryption, the second is decryption, and the third is keying. Keying is discussed last because encryption and decryption uses the golden key schedule from FIPS 197.

AES encryption

TODO

AES decryption

TODO

AES key scheduling

TODO

Chapter 5. Secure Hash Standard

SHA is the Secure Hash Standard. SHA is specified in FIPS 180-4, Secure Hash Standard (SHS) [<https://nvlpubs.nist.gov/nistpubs/fips/nist.fips.180-4.pdf>]. You should read the standard if you are not familiar with the hash family.

Sigma functions

POWER8 provides a `sigma` instruction to accelerate SHA calculations. The instruction takes two integer arguments and the constants are used to select among `Sigma0`, `Sigma1`, `sigma0` and `sigma1`.

Ch function

POWER8 provides the `vsel` instruction and it is SHA's `Ch` function. The implementation for the 32x4 arrangement is shown below. The code is the same for the 64x2 arrangement, but the function takes `uint64x2_p8` arguments. The important piece of information is `x` used as the selector.

```
uint32x4_p8
VectorCh(uint32x4_p8 x, uint32x4_p8 y, uint32x4_p8 z)
{
    return vec_sel(z, y, x);
}
```

Maj function

POWER8 provides the `vsel` instruction and it can be used for SHA's `Maj` function. The implementation for the 32x4 arrangement is shown below. The code is the same for the 64x2 arrangement, but the function takes `uint64x2_p8` arguments. The important piece of information is `x^y` used as the selector.

```
uint32x4_p8
VectorCh(uint32x4_p8 x, uint32x4_p8 y, uint32x4_p8 z)
{
    return vec_sel(y, z, vec_xor(x, y));
}
```

SHA-256

TODO

SHA-512

TODO

Chapter 6. Polynomial multiplication

The chapter of the document should discuss polynomial multiplication used with CRC codes and the GCM mode of operation for AES. However we have no experience with polynomial multiplication. Please refer to GitHub CRC32/vpmsum [<https://github.com/antonblanchard/crc32-vpmsum>].

CRC-32 and CRC-32C

No content.

GCM mode

No content.

Chapter 7. Assembly language

TODO.

Cryptogams

Cryptogams [<https://www.openssl.org/~appro/cryptogams/>] is Andy Polyakov's incubator to develop assembly language routines for OpenSSL. Andy dual licenses his implementations, so a more permissive license is available for the assembly language source code. This chapter will show you how to build Andy's software.

Chapter 8. References

The following is a list of references we are aware of.

Cryptogams

The following is a list of Cryptogams references.

- CRYPTOGRAMS: low-level cryptographic primitives collection [<https://www.openssl.org/~appro/cryptogams/>]

GitHub

The following is a list of GitHub references.

- AES Intrinsic [<https://github.com/noloader/AES-Intrinsic>]
- SHA Intrinsic [<https://github.com/noloader/SHA-Intrinsic>]
- CRC32/vpmsum [<https://github.com/antonblanchard/crc32-vpmsum>]
- sha2-le [<https://github.com/PPC64/sha2-le>]

IBM

The following is a list of IBM references.

- Recommended debug, compiler, and linker settings for Power processor tuning [<https://www.ibm.com/support/knowledgecenter/en/linuxonibm/liaal/iplsdkrecbldset.htm>]
- POWER8 in-core cryptography [<https://www.ibm.com/developerworks/library/se-power8-in-core-cryptography/index.html>]

NIST

The following is a list of NIST references.

- FIPS 197, Advanced Encryption Standard (AES) [<https://nvlpubs.nist.gov/nistpubs/fips/nist.fips.197.pdf>]
- FIPS 180-4, Secure Hash Standard (SHS) [<https://nvlpubs.nist.gov/nistpubs/fips/nist.fips.180-4.pdf>]

Stack Exchange

The following is a list of Stack Exchange references.

- Detect Power8 in-core crypto through getauxval? [<https://stackoverflow.com/q/46144668/608639>]

Index

C

Cryptogams, 14

S

SHA, 12

- Ch function, 12

- Maj function, 12

- SHA-256, 12

- SHA-512, 12

- Sigma functions, 12